

Kevin Farinango C.

Senior Data Scientist | Generative AI & ML Specialist | PhD(c) Researcher

Email: kfarinango@protonmail.com | Phone: +593 995392292 | Location: Quito, Ecuador (Open to Remote)

Portfolio: asuskf.github.io | LinkedIn: linkedin.com/in/david-farinango | GitHub: github.com/Asuskf

PROFESSIONAL SUMMARY

Senior Data Scientist and PhD(c) Researcher with 7+ years of experience in Machine Learning, NLP, and Generative AI (LLMs, RAG). Delivered scalable AI solutions through projects for global clients including Google and Microsoft, and as in-house Data Scientist at Finerio (Fintech). Expertise in cloud platforms (AWS, Azure, GCP), LLM cost optimization, and end-to-end ML pipeline design. Key achievements include a 99.86% reduction in processing time and cost using embeddings and mathematical modeling, and an 80% improvement in process efficiency through AI automation. Active researcher with 2 published papers and 2 papers under review. International speaker on Generative AI, LLMs, and Data Science.

PROFESSIONAL EXPERIENCE

Senior Data Scientist / Data Architect — Finerio (Fintech)

Feb 2024 – Feb 2026

- Spearheaded automation and optimization of data processing pipelines using Python, embeddings, and mathematical modeling — achieving a 99.86% reduction in processing time and operational cost.
- Designed and implemented scalable data lake architectures on AWS, supporting the full data lifecycle and enabling efficient processing of millions of financial transactions.
- Enhanced prompt engineering and RAG methodologies for LLM-based agents using LangChain and LangGraph, significantly improving response accuracy and contextual relevance.
- Developed advanced custom classification models leveraging embeddings and Mixture of Experts (MoE) architectures.
- Reduced LLM operational costs through strategic optimization of prompt design, model selection, and usage patterns.
- Built and deployed ML solutions on AWS SageMaker, AWS Bedrock, and Google Cloud Vertex AI.

Data Engineering Consultant — PrediQT

Mar 2023 – Jan 2024

- Led development of financial models using Python and R, increasing data processing efficiency by 30%.
- Presented actionable insights to stakeholders, contributing directly to improved financial planning and decision-making.

Data Scientist & Data Engineer — Nodel LLC

Mar 2022 – Jan 2023

- Delivered ML projects for clients including Google and British American Tobacco, optimizing ML workflows in Azure and GCP.
- Researched and deployed NLP models (BERT, Transfer Learning) improving document classification accuracy by 20%.

Senior Data Science Consultant (Independent) — Asuskf

Apr 2017 – Present

- Designed ML models (NLP, time-series) for finance and automotive sectors, improving prediction accuracy by 10%.
- Optimized Big Data infrastructure in GCP and Azure, reducing query times by 40%.
- Built ASUSKF Investing AI: autonomous decision-making engine using LLMs for real-time sentiment analysis and technical signals for market strategy execution.
- Developed Gector Evolution: advanced NLP pipeline for high-performance text cleaning and grammatical error correction, optimizing the training pipeline for faster execution.
- Key projects: sentiment analysis (Finance), signal processing for epidemiology (Mexico), time-series models for Electric Vehicle data (France).

TECHNICAL SKILLS

- **Machine Learning & AI:** Supervised/Unsupervised Learning, Deep Learning, NLP, Generative AI, LLMs, RAG, Transfer Learning, Bayesian Statistics, Time Series Analysis, Mixture of Experts (MoE)
- **GenAI Frameworks:** LangChain, LangGraph, LlamaIndex, Prompt Engineering, LLM Fine-tuning, Embeddings, Vector Databases, Annoy, HNSW
- **ML Libraries:** TensorFlow, PyTorch, Scikit-learn, Hugging Face, Pandas, NumPy, Keras
- **MLOps & Deployment:** AWS SageMaker, AWS Bedrock, Google Vertex AI, Azure ML, MLflow, Docker, FastAPI
- **Cloud & Data Engineering:** AWS, Azure, GCP, Databricks, Apache Spark, MongoDB, Redis, dbt, ETL Pipelines, Data Lake Architecture
- **Programming:** Python, R, SQL, DAX, Bash

KEY PROJECTS & OPEN SOURCE

ASUSKF Investing AI

Autonomous investment engine using LLMs for real-time market sentiment analysis and signal processing. Live public portfolio on eToro. **Tags:** LLMs, FinTech, Trading Automation

Gector Evolution

Advanced fork of the Gector architecture for high-performance text cleaning and grammatical error correction; modernized deprecated libraries and heavily optimized the training pipeline. **Tags:** NLP, Pipeline Optimization, Text Cleaning

- GitHub: github.com/Asuskf/gector-improved-training

SCIENTIFIC PUBLICATIONS

Published — Similarity Search in Diabetes Data (BRFSS)

Biomedical embeddings and Annoy trees for population health survey data. [View publication](#)

Published — Efficient Similarity Search with HNSW

[View publication](#)

Under Review — Graph-Oriented Approaches for Healthcare Fraud Detection

Under Review — Graphs Oriented to Stock Market Analysis

INTERNATIONAL SPEAKER & TALKS

- Data Governance (DAMA): Frameworks for effective governance and DataOps integration — YouTube (EN)
- Bayesian Statistics & ML: Probabilistic forecasting models for complex marketing data — YouTube (ES)
- Macroeconomics in Python: Automating financial analysis and GDP modeling via code — YouTube (ES)
- Exploring open-source Generative AI in diabetes data analysis — Escuela Politecnica Nacional, Ecuador (2024)
- Optimizing company costs using Generative AI — Escuela Kruger (2024)
- Introduction to LangChain — R4T Peru (2024)
- Introduction to Computational Neuroscience — Universidad Tecnologica Equinoccial (2024)

EDUCATION

PhD in Applied Computing

Universidad Americana de Europa, Mexico

2023 – Present

MSc. in Big Data and Data Science

2021 – 2022

UNIR, Spain

MSc. in Analysis and Visualization of Massive Data

2021 – 2022

UNIR, Mexico

Bachelor's in Systems Engineering

2015 – 2020

Universidad Tecnológica Israel, Ecuador

CERTIFICATIONS

- Digital Project Management — MIU City University Miami
- PCAP — Python Institute
- Linux Essentials — Linux Professional Institute (LPI)
- AI for Medical Diagnosis — MINTEL
- TensorFlow Specialization — Google Cloud

KEY ACHIEVEMENTS

- 99.86% reduction in AI processing time and cost using embedding, RAG, and LLM optimization (Finerio).
- 80% improvement in process efficiency through AI automation.
- 97% reduction in POC time and cost via data cleaning and clustering techniques.
- Developed an automatic prompt generation system leveraging GPT-Prompt-Engineer and Anthropic APIs.

LANGUAGES

- Spanish: Native
- English: B2 (Professional working proficiency)
- French: A1 (Basic)